

1. Wed 11/19, 11:59 PM – Scope + raw handbook scrape

- Decide on **6–8 handbook sections** to cover (e.g., alcohol policy, housing/quiet hours, academic honesty, absences, Title IX, campus safety).
 - Scrape/download those sections from the **online Centre Student Handbook** and save each one as a **separate .txt file** in `data/` (one file per section).
 - Make a rough list of **20–30 likely student questions** that could be answered from those sections (not the gold eval set yet).
-

2. Sun 11/23, 11:59 PM – Chunking + metadata finished

- Write `src/ingest.py` and run it on the `.txt` files to create `index/chunks.jsonl`.
 - Ensure each chunk has metadata:
`{section_title, url, last_updated_if_known, chunk_id, text}` and is under **~700–900 words**.
-

3. Mon 12/1, 10:00 PM – Doc2Query questions + embeddings index

- For **every chunk**, write **3–5 student-style questions** and save them as `index/questions.jsonl` with `{question, chunk_id}`.
 - Implement `src/index.py` to:
 - Embed all generated questions with the chosen embedding model.
 - Save the vector matrix (e.g., `embeddings.npy`) and a question→chunk mapping file.
 - Confirm you can run `index.py` end-to-end without errors.
-

4. Wed 12/3, 11:59 PM – Basic REPL working end-to-end

- Implement `src/ask.py` as a CLI REPL that:
 - Embeds user queries, retrieves **top-k similar generated questions**, pulls linked chunks.
 - Returns a **1–2 sentence answer strictly from the chunk**, plus a citation: *(section title + URL)*.
 - Says **“insufficient context”** when similarities are low or no chunk clearly matches.
- Manually test with **10–15 informal questions** (not the gold set yet) and note any obvious failures.

5. Thu 12/4, 11:00 PM — Evaluation + tuning + README skeleton

- Write **10 gold questions** up front that you think real students would ask.
- Run an evaluation loop to check whether the **correct chunk is in the top-3** for each gold question and whether the surfaced text contains the exact number/phrase when needed.
- Adjust retrieval settings (top-k, similarity threshold, “insufficient context” behavior) based on results.
- Draft the structure of **README.md** (run instructions + place for the 10 gold Q/A with pass/fail marks).

6. Fri 12/5, 3:00 PM — Final polish + submission

- Final pass on **code and data folders** (`data/`, `index/`, `src/`), remove unused files, ensure scripts have clear usage instructions or `--help`.
- Finish **README.md**:
 - How to set up and run the project.
 - Your 10 gold questions, the returned answers, and pass/fail marks.
- Write/clean up the **short written report** (if separate from README) summarizing what works, what doesn't, and how you evaluated.
- Confirm everything runs from a clean environment and is ready to upload **before 4:00 PM**.