

1. Progress so far

- From the Check-In 1 plan, I have completed: (a) choosing the initial scope of ~8 high-impact handbook policies (alcohol/drugs, Good Samaritan, Greek life, hazing, conduct, sexual misconduct/Title IX, travel, visitation), (b) building the ingest pipeline to clean the .txt files and chunk them into ~58 paragraph-based chunks, and (c) standing up a first working retrieval pipeline.
- Implemented doc2query with a local LLM (qwen2.5:1.5b via Ollama): it generates 3–4 student-style questions per chunk, with a heuristic filter (word-overlap + number consistency) to drop hallucinated or off-topic questions.
- Built the vector index using mxbai-embed-large: embeddings.npy + meta.json are in place and wired to the synthetic questions.
- Implemented ask.py: REPL embeds user queries, retrieves the top-1 relevant chunk via cosine similarity, then answers strictly from that chunk using a constrained prompt; returns “Insufficient context...” when similarity is low.
- Added URL metadata per policy in chunks.jsonl so the REPL can print sources in the form “Policy Name – Student Handbook (URL)” and sanity-tested behavior on several queries (Good Samaritan, visitation/overnight guests, out-of-scope questions like unexcused absences).

2. Plan for the final week (Dec 1–5)

- Mon 12/1 (by 8:00 pm): Quick manual pass over a sample of synthetic doc2query questions to confirm quality; adjust similarity threshold (MIN_SIMILARITY) only if obviously needed. (Responsible: Gabriel)
- Tue 12/2 (by 10:00 pm): Create gold.jsonl with ~10–15 hand-written “gold” questions covering all chosen policies, each with a gold answer and expected section; run the system on all gold questions in the REPL and log outcomes (correct / incorrect / insufficient). (Responsible: Gabriel)
- Wed 12/3 (by 9:00 pm): Freeze all LLM-dependent runs (doc2query generation, embedding, and evaluation runs); commit final index/ artifacts so no further model changes occur after this point. (Responsible: Gabriel)
- Thu 12/4 (by 11:00 pm): Draft the written report and project README, including system design, chunking/doc2query/embedding choices, REPL behavior, and summarized gold-set results; revise for clarity. (Responsible: Gabriel)

- Fri 12/5 (by 3:00 pm): Final code cleanup and sanity checks, confirm all code, data, and documentation are pushed/archived and any required uploads are completed before the 4:00 pm deadline. (Responsible: Gabriel)